

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

Generative Models


Outline

- What are generative models
- Types of generative models



Generative Models vs. Discriminative Models

Discriminative models



Features Class

$$X \rightarrow Y$$

$$P(Y|X)$$

Generative models



Noise Class Features

$$\xi, Y \rightarrow X$$

$$P(X|Y)$$

Generative Models vs. Discriminative Models



Generative models




Noise Class Features

$$\xi, Y \rightarrow X$$

$$P(X|Y)$$

Generative Models

Variational Autoencoders

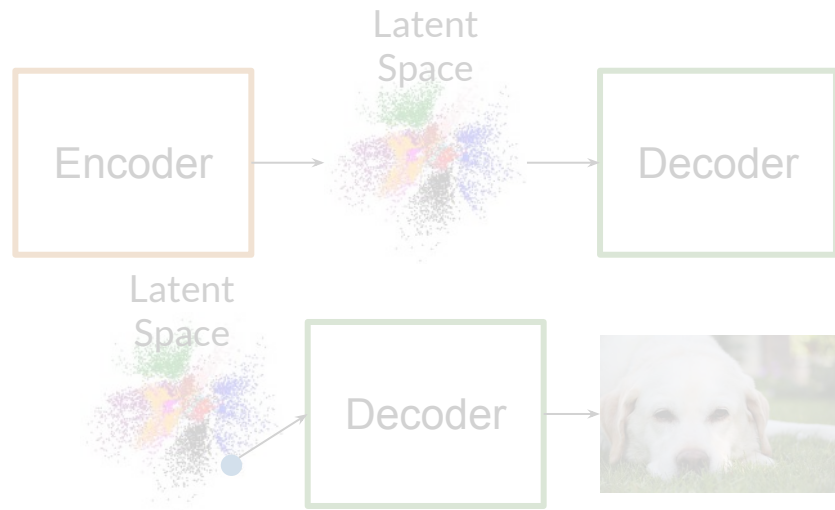


Generative Adversarial Networks


Available from: <https://arxiv.org/abs/1804.00891>

Generative Models

Variational Autoencoders




Generative Adversarial Networks



Available from: <https://arxiv.org/abs/1804.00891>

Summary

- Generative models learn to produce examples
- Discriminative models distinguish between classes
- Up next, GANs!






deeplearning.ai

Real Life GANs

Outline

- Cool applications of GANs
- Major companies using them



GANs Over Time



Ian Goodfellow
@goodfellow_ian




4.5 years of GAN progress on face generation.

arxiv.org/abs/1406.2661 arxiv.org/abs/1511.06434

arxiv.org/abs/1606.07536 arxiv.org/abs/1710.10196

arxiv.org/abs/1812.04948



GANs Over Time

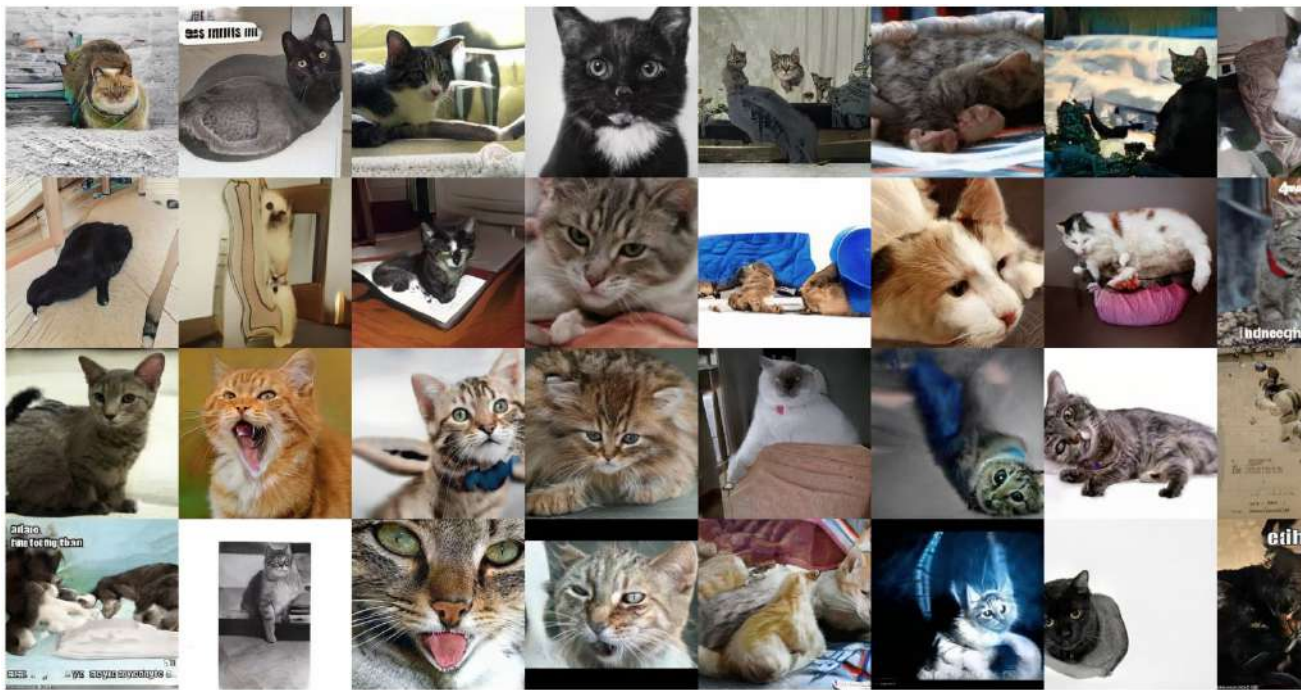


Face Generation
StyleGAN2

These people do
not exist!

Karras, Tero, et al. "Analyzing and improving the image quality of stylegan." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020.

GANs Over Time



StyleGAN2



Mimics the
distribution of the
training data

Karras, Tero, et al. "Analyzing and improving the image quality of stylegan." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020.

<https://9gag.com/gag/aWYZKWx>

GANs for Image Translation


From one domain to another

CycleGAN



Park, Taesung, et al. "Semantic image synthesis with spatially-adaptive normalization." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2019.

GANs for Image Translation



GauGAN

Doodles



Pictures

Park, Taesung, et al. "Semantic image synthesis with spatially-adaptive normalization." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2019.

GANs are Magic!



Zakharov, Egor, et al. "Few-shot adversarial learning of realistic neural talking head models." *Proceedings of the IEEE International Conference on Computer Vision*. 2019.


GANs for 3D Objects

Chair




3D-GAN


Car



Sofa



Table



Generative
Design

Wu, Jiajun, et al. "Learning a probabilistic latent space of object shapes via 3d generative-adversarial modeling." *Advances in neural information processing systems*. 2016.

Companies Using GANs



Next-gen
Photoshop



Text
Generation



Data
Augmentation





Image Filters



Super-resolution

Summary

- GANs' performance is rapidly improving
- Huge opportunity to work in this space!
- Major companies are using them





deeplearning.ai

Intuition Behind GANs


Outline

- The goal of the generator and the discriminator
- The competition between them




Generative Adversarial Network

Generator learns to make *fakes*
that look *real*



Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Generator learns to make *fakes*
that look **real**

Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Generator learns to make *fakes*
that look *real*



Doesn't know how
it should look



Generative Adversarial Network

Discriminator learns to distinguish
real from *fake*



The Game Is On!



5% Real




40% Real



80% Real




The Game is On!



The Game Is On!



30% Real



60% Real




95% Real

The Game Is On!




30% Real



60% Real



Wrong!




95% Real



Won't fool me
again



The Game Is On!



Summary

- The **generator's** goal is to fool the discriminator
- The **discriminator's** goal is to distinguish between real and fake
- They learn from the competition with each other
- At the end, *fakes* look **real**






deeplearning.ai

Discriminator

Outline

- Review of classifiers
- The role of classifiers in terms of probability
- Discriminator



Classifiers

Distinguish between different classes



Turtle

Bird


Cat

Dog


Fish

Classifiers


Distinguish between different classes



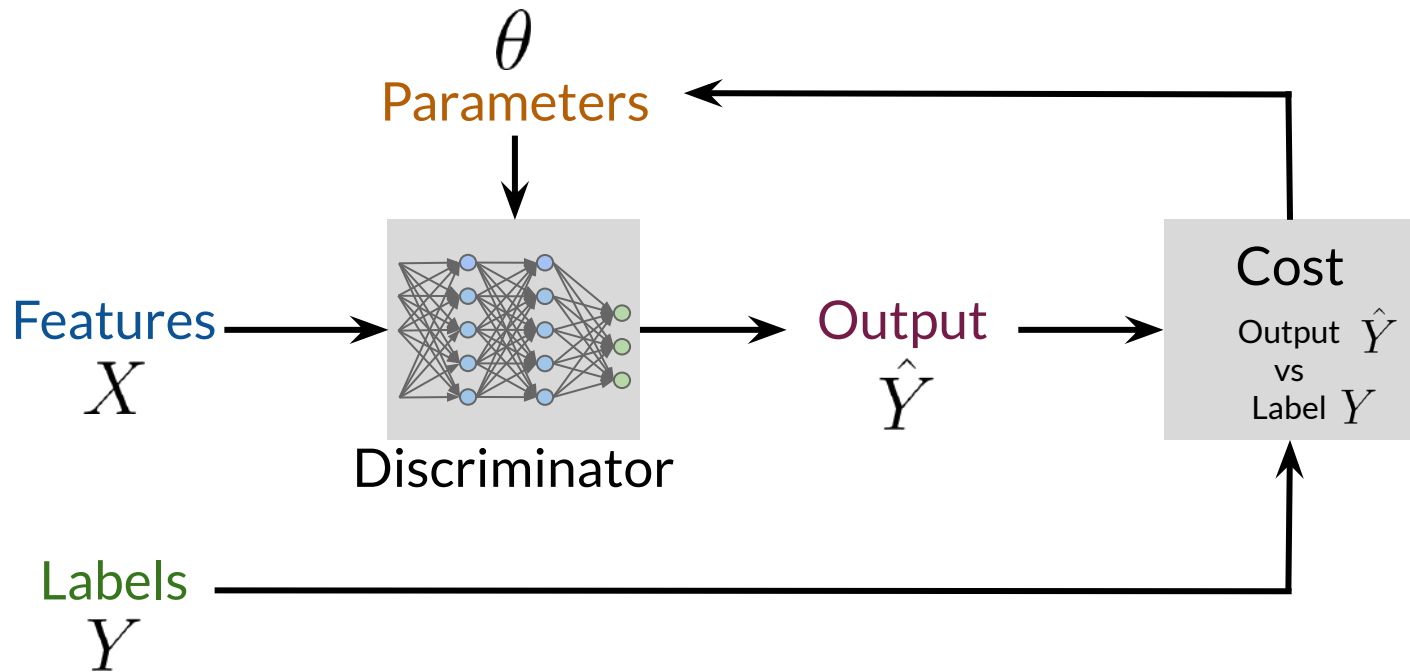
Neural Networks



Neural Networks




Classifiers (training)




Classifiers



$$P\left(\begin{array}{c} \text{Turtle} \\ \text{Bird} \\ \text{Cat} \\ \text{Dog} \\ \text{Fish} \end{array} \mid \text{Image} \right)$$





Classifiers



Discriminator



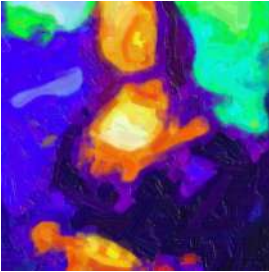
Discriminator



Discriminator

$$P\left(\begin{array}{c} \text{Fake} \\ \text{Class} \end{array} \mid \begin{array}{c} X \\ \text{Features} \end{array} \right)$$

Discriminator

$$P\left(\begin{array}{c} \text{Fake} \\ \text{Class} \end{array} \mid \begin{array}{c} \text{Features} \end{array}\right) = 0.85 \rightarrow \boxed{\text{Fake}}$$


Summary

- The **discriminator** is a classifier
- It learns the probability of class Y (**real** or *fake*) given features X
- The probabilities are the feedback for the **generator**





deeplearning.ai

Generator

Outline

- What the generator does
- How it improves its performance
- Generator in terms of probability



Generator

Turtle

Generates examples of the class

Bird


Cat

Dog


Fish



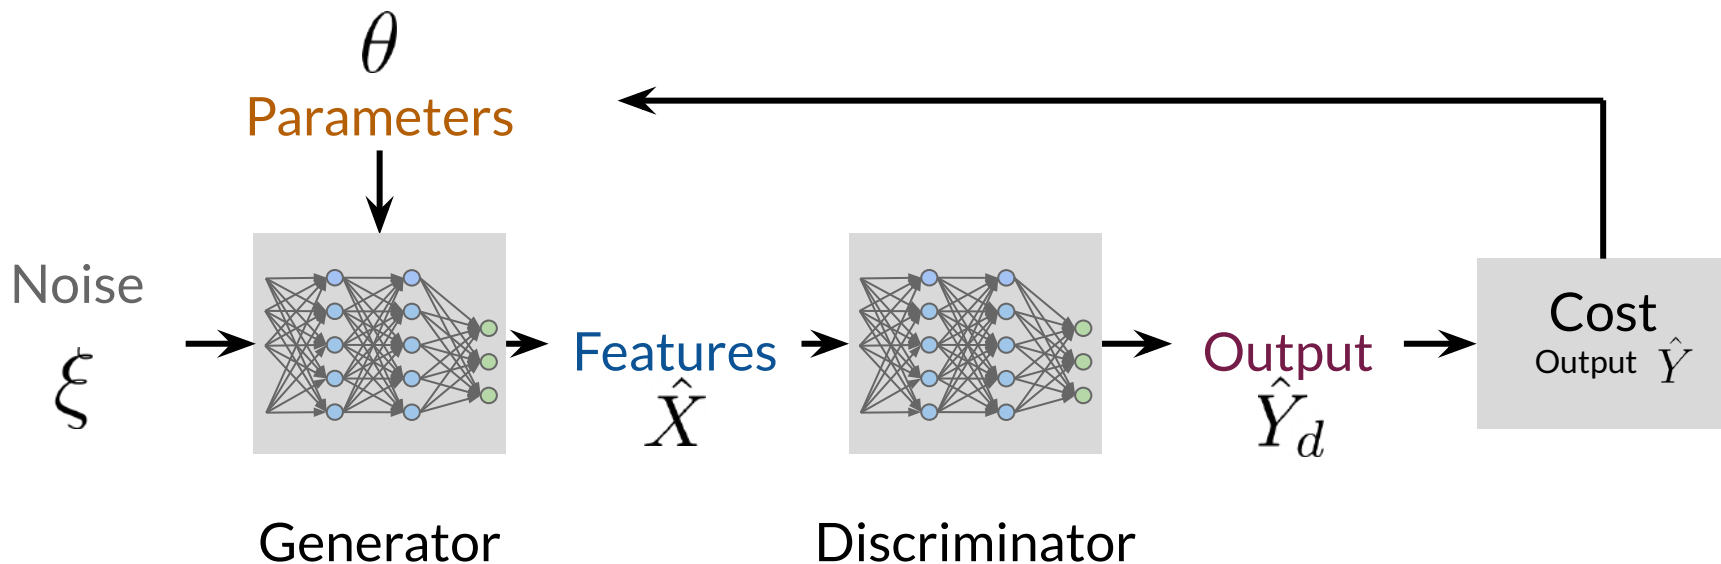
Neural Networks




Neural Networks



Generator: Learning



Sampling



Generator

$$P(\text{Image} \mid \text{Class})$$

Turtle


Bird

Cat

Dog

Fish

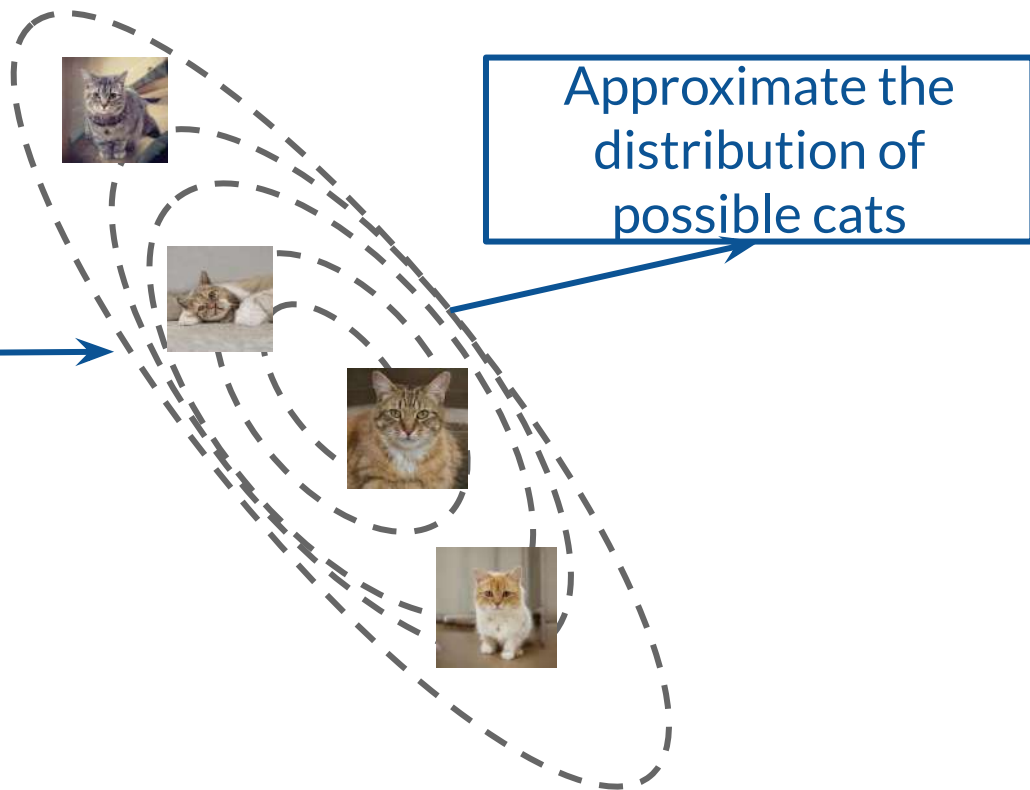
Generator



Generator

$P(X)$

Features



Images available from: <http://thesecatsdonotexist.com/>

Summary

- The **generator** produces fake data
- It learns the probability of features X
- The **generator** takes as input noise (random features)





deeplearning.ai

BCE Cost Function

Outline

- Binary Cross Entropy (BCE) Loss equation by parts
- How it looks graphically



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Prediction

Label

Features

Parameters

Average loss of the whole batch

The diagram illustrates the components of the Binary Cross-Entropy (BCE) cost function. The formula is shown as $J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$. Annotations include: a circle around the summation term with an arrow pointing to a box labeled 'Average loss of the whole batch'; an arrow from $h(x^{(i)}, \theta)$ pointing to 'Prediction'; an arrow from $y^{(i)}$ pointing to 'Label'; an arrow from $x^{(i)}$ pointing to 'Features'; and an arrow from θ pointing to 'Parameters'.

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

| $y^{(i)}$ | $h(x^{(i)}, \theta)$ | $y^{(i)} \log h(x^{(i)}, \theta)$ |
|-----------|----------------------|-----------------------------------|
| 0 | any | 0 |
| 1 | 0.99 | ~ 0 |
| 1 | ~ 0 | -inf |

Relevant when
the label is 1

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

| $y^{(i)}$ | $h(x^{(i)}, \theta)$ | $(1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))$ |
|-----------|----------------------|--|
| 1 | any | 0 |
| 0 | 0.01 | ~0 |
| 0 | ~1 | -inf |

Relevant when
the label is 0


BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Ensures that the cost is always greater or equal to 0



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



Summary

- The BCE cost function has two parts (one relevant for each class)
- Close to zero when the label and the prediction are similar
- Approaches infinity when the label and the prediction are different






deeplearning.ai


Putting It All Together

Outline

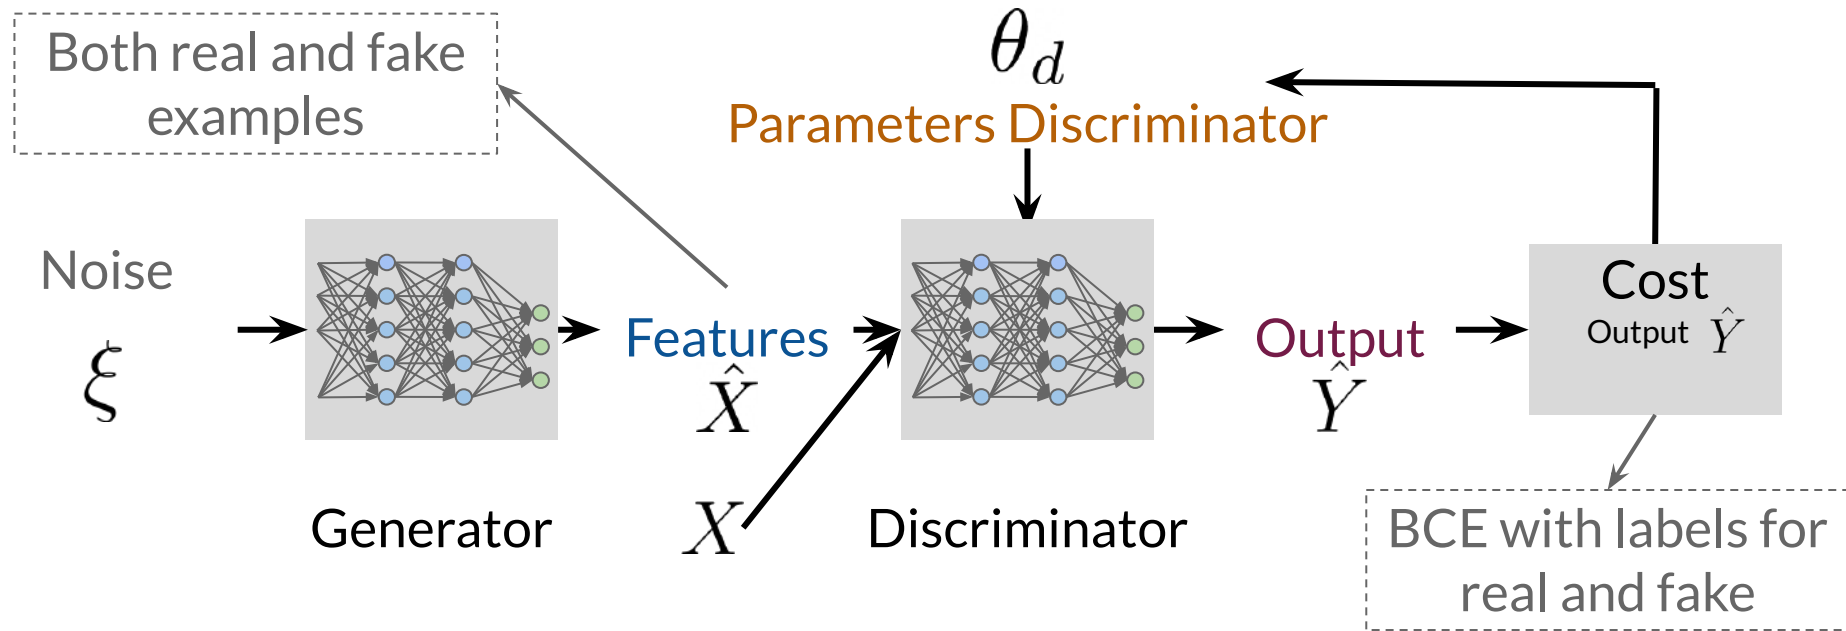
- How the whole architecture looks
- How to train GANs



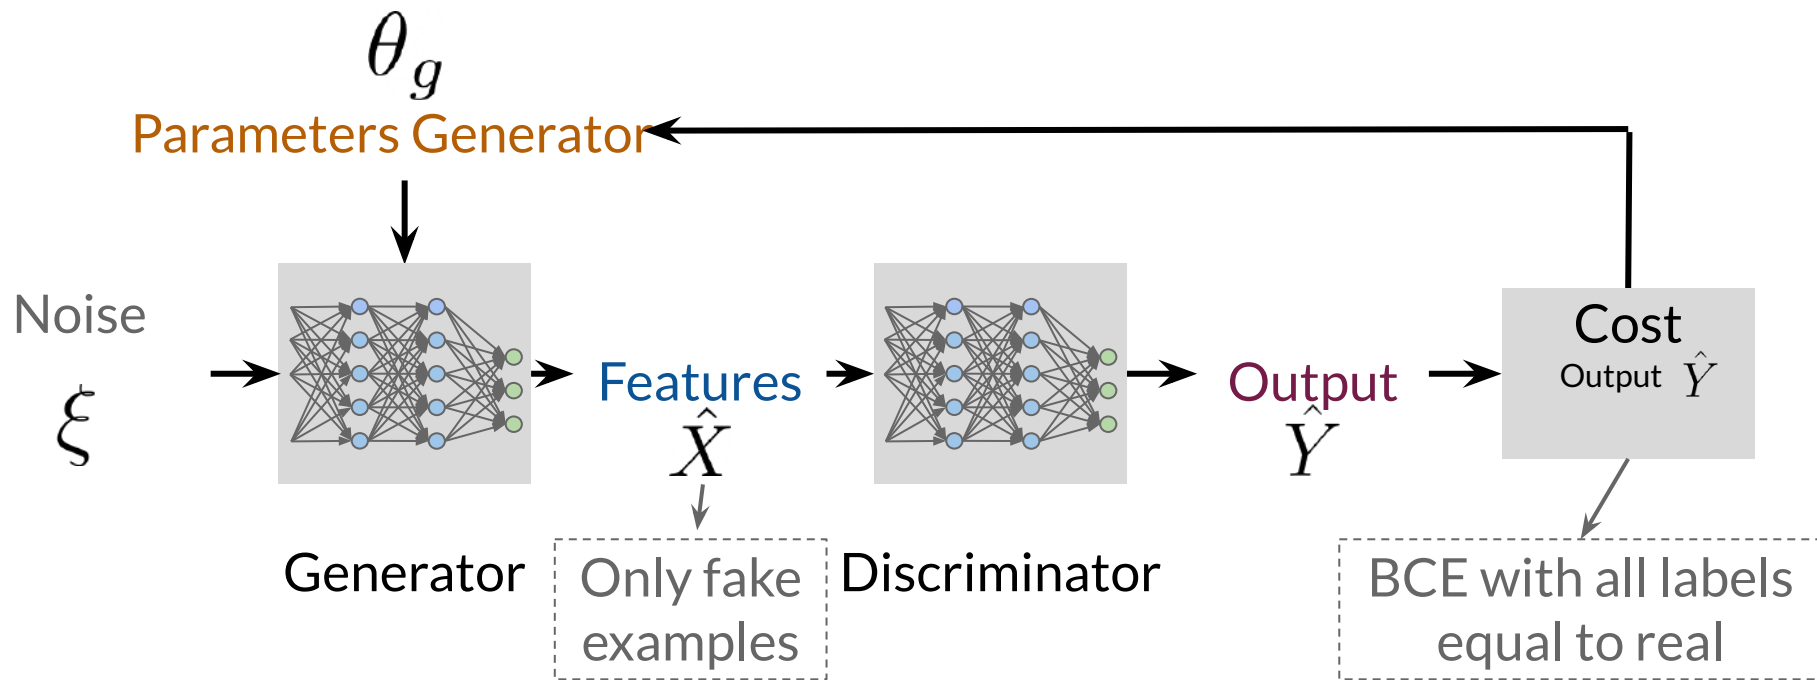
GANs Model



Training GANs: Discriminator



Training GANs: Generator




Training GANs



Summary

- GANs train in an alternating fashion
- The two models should always be at a similar “skill” level





deeplearning.ai

Intro to PyTorch (Optional)

Outline

- Comparison with TensorFlow
- Defining Models
- Training



PyTorch vs TensorFlow

PyTorch

Imperative, computations on the go

```
A, B = 1, 2
C = A + B
print(C)
```

3

Dynamic Computational Graphs

TensorFlow

Symbolic, first define and then compile

```
C = A + B
f = compile(C)
print(f(A = 1, B = 2))
```

3

Static Computational Graphs

Tensorflow > 2.0 moves toward PyTorch by including **Eager Execution**

PyTorch vs TensorFlow

PyTorch

TensorFlow

Currently very similar frameworks!

Tensorflow > 2.0 moves toward PyTorch by
including Eager Execution

Defining Models in PyTorch

```
import torch
from torch import nn
```

→ Custom layers for DL

```
class LogisticRegression(nn.Module):
    def __init__(self, in):
        super().__init__()
        self.log_reg = nn.Sequential(
            nn.Linear(in, 1),
            nn.Sigmoid()
        )
    def forward(self, x):
        return self.log_reg(x)
```

Define the model as a class

Initialization method with parameters

Definition of the architecture

Forward computation of the model
with inputs x

Training Models In PyTorch

```
model = LogisticRegression(16)
```

Initialization of the model

```
criterion = nn.BCELoss()
```

Cost function

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

Optimizer

```
for t in range(n_epochs):
```

Training loop for number of epochs

```
    y_pred = model(x)
```

```
    loss = criterion(y_pred, y)
```

Forward propagation

```
    optimizer.zero_grad()
```

```
    loss.backward()
```

```
    optimizer.step()
```

Optimization step

Summary

- PyTorch makes computations on the run
- Dynamic computational graphs in Pytorch
- Just another framework, and similar to Tensorflow!

